

Machine Learning Assignment

Logitboost

1 a

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, roc_auc_score
from sklearn.preprocessing import StandardScaler
```

```
In [6]: def logistic_regression_scaled(X_train, y_train, X_test, y_test, k_values, learning_rate=0.1, regularization_strength=0.01):

    scaler = StandardScaler()
    X_train_scaled = scaler.fit_transform(X_train)
    X_test_scaled = scaler.transform(X_test)

    ones_train = np.ones(X_train_scaled.shape[0])
    ones_test = np.ones(X_test_scaled.shape[0])
    train_data = np.insert(X_train_scaled, 0, ones_train, axis=1)
    test_data = np.insert(X_test_scaled, 0, ones_test, axis=1)
    N_train = train_data.shape[0]
    M_train = train_data.shape[1]

    loss_list = []
    error_train_list = []
    error_test_list = []
    results_list = []

    for k in k_values:
        X_train_copy = train_data.copy()
        X_test_copy = test_data.copy()
        beta = np.zeros(M_train)

        for _ in range(0, k):
            l_pred = np.dot(X_train_copy, beta)
            prob = 1.0 / (1.0 + np.exp(-2 * l_pred))
            weights = prob * (1.0 - prob)
            resid = 0.5 * (y_train + 1) - prob
            resid[weights == 0] = 0
            resid[weights != 0] = resid[weights != 0]
            / weights[weights != 0]
            coeff = np.zeros((2, M_train - 1))
            new_loss = np.zeros(M_train - 1)

            for j in range(0, M_train - 1):
                X_j = X_train_copy[:, j + 1]
                sum_w = np.sum(weights)
                sum_w_X_j = np.sum(weights * X_j)
                sum_w_X_j_squared = np.sum(weights * X_j ** 2)
                sum_w_residual = np.sum(weights * resid)
                sum_w_X_j_residual = np.sum(weights * X_j * resid)

                if (sum_w * sum_w_X_j_squared - sum_w_X_j ** 2) == 0:
                    beta_j = np.array([sum_w_residual / sum_w, 0])
```

```

    else:
        beta_j = np.array([sum_w_X_j_squared *
                           sum_w_residual - sum_w_X_j *
                           sum_w_X_j_residual, sum_w *
                           sum_w_X_j_residual - sum_w_X_j
                           * sum_w_residual]) /
        (sum_w * sum_w_X_j_squared - sum_w_X_j ** 2)

        l_pred_j = l_pred + 0.5 * (beta_j[0] + beta_j[1] * X_j)
        loss_j = np.sum(np.log(1 + np.exp(-2 * y_train * l_pred_j)))
        coeff[:, j] = beta_j
        new_loss[j] = loss_j

    min_loss_index = np.argmin(new_loss)
    beta[0] = beta[0] + learning_rate *
    (0.5 * coeff[0, min_loss_index] - regularization_strength * beta[0])
    beta[min_loss_index + 1] = beta[min_loss_index + 1]
    + learning_rate * (0.5 * coeff[1][min_loss_index] - regularization_strength
                      * beta[min_loss_index + 1])

    if k == k_values[-1]:
        loss_list.append(new_loss[min_loss_index])

```

```

l_pred_train = np.dot(X_train_copy, beta)
pred_train = np.where(l_pred_train > 0.0, 1, -1)
error_train = 1 - np.mean(pred_train == y_train)
error_train_list.append(error_train)

```

```

l_pred_test = np.dot(X_test_copy, beta)
pred_test = np.where(l_pred_test > 0.0, 1, -1)
error_test = 1 - np.mean(pred_test == y_test)
error_test_list.append(error_test)
results_list.append({"k": k, "Training Errors": error_train,
                    "TestErrors": error_test})

```

```

if k == 300:
    p_train_300 = l_pred_train
    pred_train_300 = pred_train

    p_test_300 = l_pred_test
    pred_train_300 = pred_test

    # Calculate ROC curve and ROC area for training
    train_fpr_300, train_tpr_300, _ = roc_curve(y_train, p_train_300)
    train_auc_300 = roc_auc_score(y_train, p_train_300)

    # Calculate ROC curve and ROC area for testing
    test_fpr_300, test_tpr_300, _ = roc_curve(y_test, p_test_300)
    test_auc_300 = roc_auc_score(y_test, p_test_300)

    # Plot ROC curves
    plt.figure(figsize=(10, 7))
    plt.plot(train_fpr_300, train_tpr_300,
             label=f'Training AUC (k=300): {train_auc_300:.2f}')
    plt.plot(test_fpr_300, test_tpr_300,
             label=f'Testing AUC (k=300): {test_auc_300:.2f}')
    plt.plot([0, 1], [0, 1], 'k--') # Dashed diagonal
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC Curve for k=300')
    plt.legend(loc="lower right")
    plt.show()

```

```

results_df = pd.concat([pd.DataFrame(result, index=[0])

```

```

        for result in results_list], ignore_index=True)

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(range(1, 601), loss_list)
plt.xlabel('Iteration')
plt.ylabel('Training Loss')
plt.show()

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(k_values, error_train_list, label='Training')
plt.plot(k_values, error_test_list, label='Testing')
plt.xlabel('K Values')
plt.ylabel('Misclassification Error for Train and Test Data')
plt.legend()
plt.show()

return results_df

# Load data
X_train_data = np.loadtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_6/arcene/arcene_train.data")
y_train_labels = np.loadtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_6/arcene/arcene_train.labels")
X_test_data = np.loadtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_6/arcene/arcene_valid.data")
y_test_labels = np.loadtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_6/arcene/arcene_valid.labels")

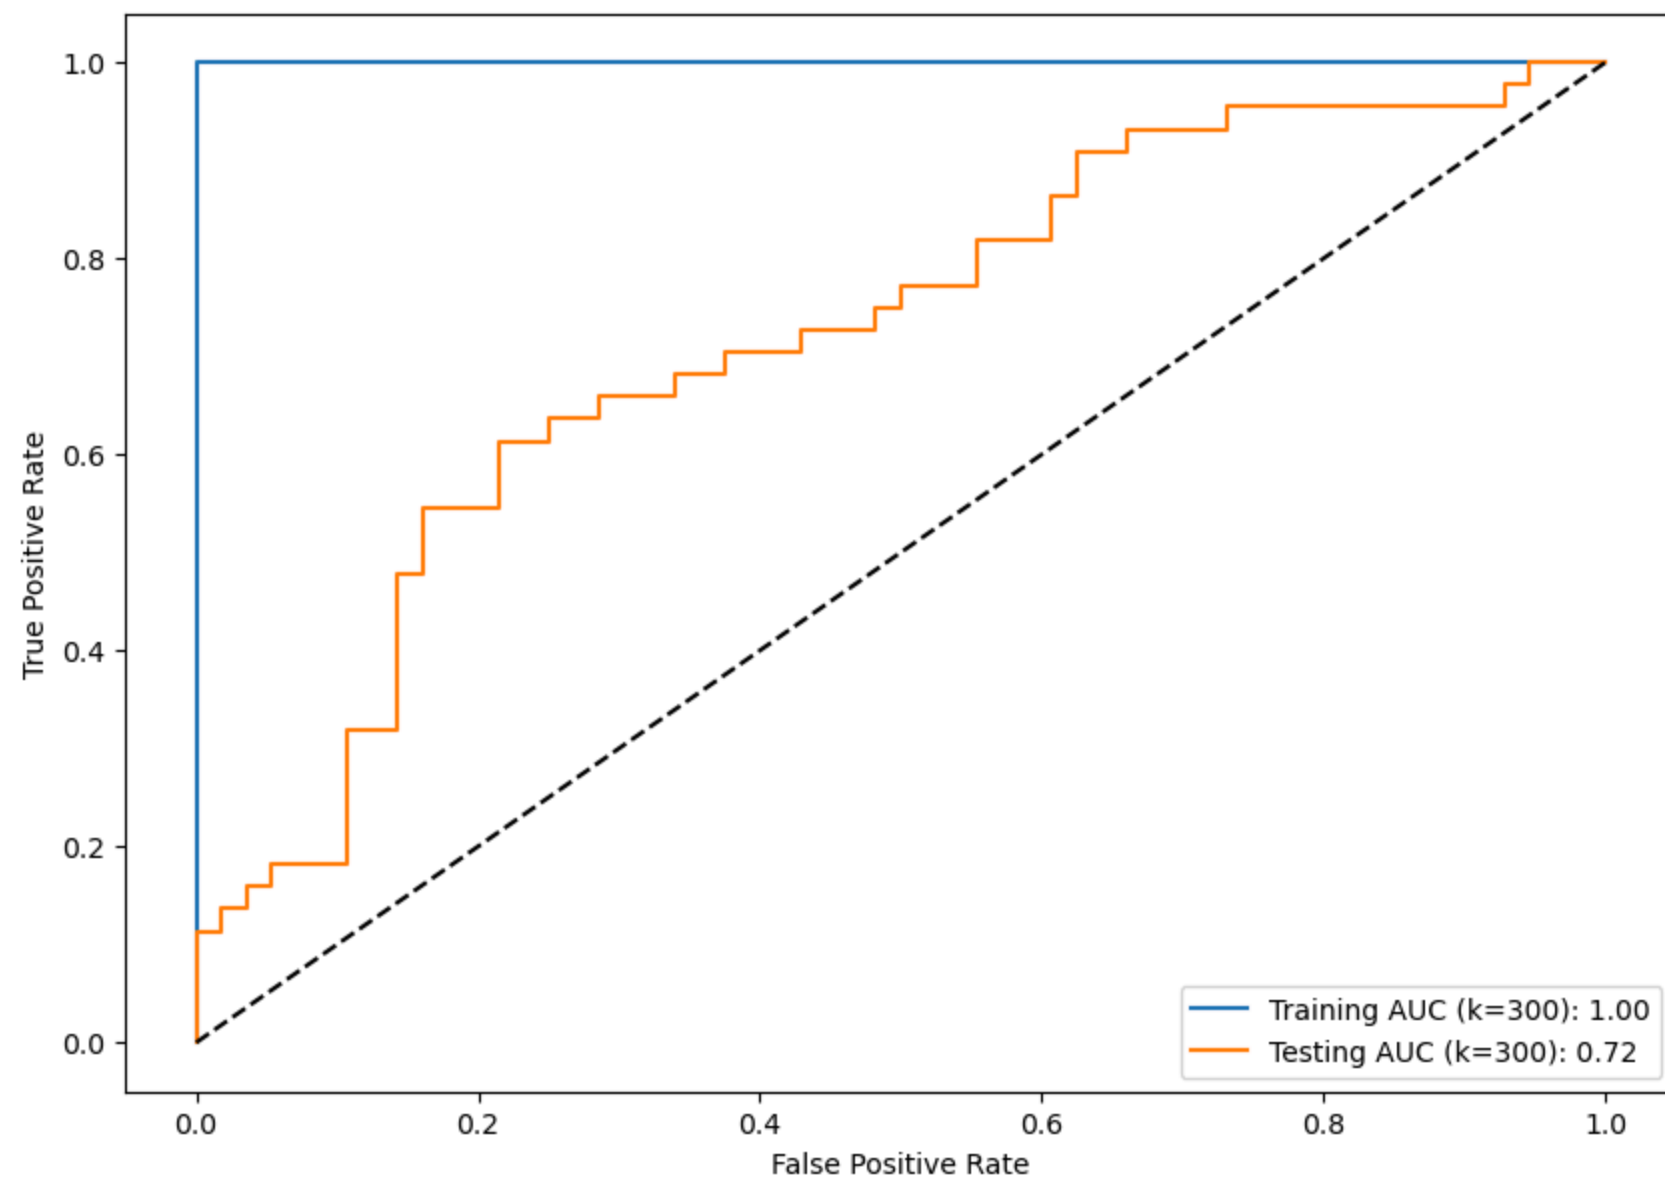
# Define values of k
k_values = [10, 30, 100, 300, 600]



# Call the logistic regression function with feature scaling
results_df_scaled = logistic_regression_scaled(X_train_data, y_train_labels,
                                              X_test_data, y_test_labels, k_values)

print("Misclassification Error for Train and Test Data:")
print(results_df_scaled)

```

ROC Curve for k=300





Misclassification Error for Train and Test Data:

	k	Training Errors	TestErrors
0	10	0.20	0.34
1	30	0.06	0.35
2	100	0.01	0.33
3	300	0.00	0.32
4	600	0.00	0.32

1b

```
In [5]: def logistic_regression_logit(X_train, y_train, X_test, y_test, k_values):
    ones_train = np.ones(X_train.shape[0])
    ones_test = np.ones(X_test.shape[0])
    train_data = np.insert(X_train, 0, ones_train, axis=1)
    test_data = np.insert(X_test, 0, ones_test, axis=1)
    N_train = train_data.shape[0]
    M_train = train_data.shape[1]

    loss_list = []
    error_train_list = []
    error_test_list = []
    results_list = []
    p_train_300 = None
    pred_train_300 = None

    for k in k_values:
        X_train_copy = train_data.copy()
        X_test_copy = test_data.copy()
        beta = np.zeros(M_train)

        for _ in range(0, k):
            l_pred = np.dot(X_train_copy, beta)
            probability = 1.0 / (1.0 + np.exp(-2 * l_pred))
            weights = probability * (1.0 - probability)
            residual = 0.5 * (y_train + 1) - probability
            residual[weights == 0] = 0
            residual[weights != 0] = residual[weights != 0]
            / weights[weights != 0]
            coeff = np.zeros((2, M_train - 1))
            new_loss = np.zeros(M_train - 1)

            for j in range(0, M_train - 1):
                X_j = X_train_copy[:, j + 1]
                sum_w = np.sum(weights)
                sum_w_X_j = np.sum(weights * X_j)
                sum_w_X_j_squared = np.sum(weights * X_j ** 2)
                sum_w_residual = np.sum(weights * residual)
                sum_w_X_j_residual = np.sum(weights * X_j * residual)

                if (sum_w * sum_w_X_j_squared - sum_w_X_j ** 2) == 0:
                    beta_j = np.array([sum_w_residual / sum_w, 0])
                else:
                    beta_j = np.array([sum_w_X_j_squared *
                                         sum_w_residual - sum_w_X_j * sum_w_X_j_residual, sum_w *
                                         sum_w_X_j_residual - sum_w_X_j * sum_w_residual]) /
                                (sum_w * sum_w_X_j_squared - sum_w_X_j ** 2)

                l_pred_j = l_pred + 0.5 * (beta_j[0] + beta_j[1] * X_j)
                loss_j = np.sum(np.log(1 + np.exp(-2 * y_train * l_pred_j)))
                coeff[:, j] = beta_j
                new_loss[j] = loss_j

            min_loss_index = np.argmin(new_loss)
            beta[0] = beta[0] + 0.5 * coeff[0, min_loss_index]
            beta[min_loss_index + 1] = beta[min_loss_index + 1]
            + 0.5 * coeff[1][min_loss_index]

        if k == k_values[-1]:
            loss_list.append(new_loss[min_loss_index])
```

```

l_pred_train = np.dot(X_train_copy, beta)
pred_train = np.where(l_pred_train > 0.0, 1, -1)
error_train = 1 - np.mean(pred_train == y_train)
error_train_list.append(error_train)

l_pred_test = np.dot(X_test_copy, beta)
pred_test = np.where(l_pred_test > 0.0, 1, -1)
error_test = 1 - np.mean(pred_test == y_test)
error_test_list.append(error_test)
results_list.append({"k": k, "Training Errors": error_train,
                    "TestErrors": error_test})

if k == 300:
    p_train_300 = l_pred_train
    pred_train_300 = pred_train

    p_test_300 = l_pred_test
    pred_train_300 = pred_test

    # Calculate ROC curve and ROC area for training
    train_fpr_300, train_tpr_300, _ = roc_curve(y_train, p_train_300)
    train_auc_300 = roc_auc_score(y_train, p_train_300)

    # Calculate ROC curve and ROC area for testing
    test_fpr_300, test_tpr_300, _ = roc_curve(y_test, p_test_300)
    test_auc_300 = roc_auc_score(y_test, p_test_300)

    # Plot ROC curves
    plt.figure(figsize=(10, 7))
    plt.plot(train_fpr_300, train_tpr_300,
             label=f'Training AUC (k=300): {train_auc_300:.2f}')
    plt.plot(test_fpr_300, test_tpr_300,
             label=f'Testing AUC (k=300): {test_auc_300:.2f}')
    plt.plot([0, 1], [0, 1], 'k--') # Dashed diagonal
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC Curve for k=300')
    plt.legend(loc="lower right")
    plt.show()

results_df = pd.concat([pd.DataFrame(result, index=[0])
                        for result in results_list], ignore_index=True)

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(range(1, 601), loss_list)
plt.xlabel('Iteration')
plt.ylabel('Training Loss for K= 600')
plt.show()

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(k_values, error_train_list, label='Training')
plt.plot(k_values, error_test_list, label='Testing')
plt.xlabel('K Values')
plt.ylabel('Misclassification Error for Train and Test Data')
plt.legend()
plt.show()

return results_df

```

Load data

```

X_train_data = np.genfromtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_3/dexter/dexter_train.csv", delimiter=',')
y_train_labels = np.genfromtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_3/dexter/dexter_train.labels")


X_test_data = np.genfromtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_3/dexter/dexter_valid.csv", delimiter=',')
y_test_labels = np.genfromtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_3/dexter/dexter_valid.labels")



# Define values of k
k_values = [10, 30, 100, 300, 600]

# Call the logistic regression function
results_df = logistic_regression_logit(X_train_data, y_train_labels,
                                       X_test_data, y_test_labels, k_values)

print("Misclassification Error for Train and Test Data:")
print(results_df)

```





Misclassification Error for Train and Test Data:

	k	Training Errors	TestErrors
0	10	0.073333	0.106667
1	30	0.003333	0.093333
2	100	0.000000	0.090000
3	300	0.000000	0.083333
4	600	0.000000	0.090000

1C

```
In [3]: def logistic_regression_logit(X_train, y_train, X_test, y_test, k_values):
    ones_train = np.ones(X_train.shape[0])
    ones_test = np.ones(X_test.shape[0])
    train_data = np.insert(X_train, 0, ones_train, axis=1)
    test_data = np.insert(X_test, 0, ones_test, axis=1)
    N_train = train_data.shape[0]
    M_train = train_data.shape[1]

    loss_list = []
    error_train_list = []
    error_test_list = []
    results_list = []
    p_train_300 = None
    pred_train_300 = None

    for k in k_values:
        X_train_copy = train_data.copy()
        X_test_copy = test_data.copy()
        beta = np.zeros(M_train)

        for _ in range(0, k):
            l_pred = np.dot(X_train_copy, beta)
            probability = 1.0 / (1.0 + np.exp(-2 * l_pred))
            weights = probability * (1.0 - probability)
            residual = 0.5 * (y_train + 1) - probability
            residual[weights == 0] = 0
            residual[weights != 0] = residual[weights != 0]
            / weights[weights != 0]
            coeff = np.zeros((2, M_train - 1))
            new_loss = np.zeros(M_train - 1)

            for j in range(0, M_train - 1):
                X_j = X_train_copy[:, j + 1]
                sum_w = np.sum(weights)
                sum_w_X_j = np.sum(weights * X_j)
                sum_w_X_j_squared = np.sum(weights * X_j ** 2)
                sum_w_residual = np.sum(weights * residual)
                sum_w_X_j_residual = np.sum(weights * X_j * residual)

                if (sum_w * sum_w_X_j_squared - sum_w_X_j ** 2) == 0:
                    beta_j = np.array([sum_w_residual / sum_w, 0])
                else:
                    beta_j = np.array([sum_w_X_j_squared *
                                         sum_w_residual - sum_w_X_j * sum_w_X_j_residual,
                                         sum_w * sum_w_X_j_residual - sum_w_X_j * sum_w_residual]) /
                                (sum_w * sum_w_X_j_squared - sum_w_X_j ** 2)

                l_pred_j = l_pred + 0.5 * (beta_j[0] + beta_j[1] * X_j)
                loss_j = np.sum(np.log(1 + np.exp(-2 * y_train * l_pred_j)))
                coeff[:, j] = beta_j
                new_loss[j] = loss_j

            min_loss_index = np.argmin(new_loss)
            beta[0] = beta[0] + 0.5 * coeff[0, min_loss_index]
            beta[min_loss_index + 1] = beta[min_loss_index + 1] +
            0.5 * coeff[1][min_loss_index]

        if k == k_values[-1]:
            loss_list.append(new_loss[min_loss_index])
```

```

l_pred_train = np.dot(X_train_copy, beta)
pred_train = np.where(l_pred_train > 0.0, 1, -1)
error_train = 1 - np.mean(pred_train == y_train)
error_train_list.append(error_train)

l_pred_test = np.dot(X_test_copy, beta)
pred_test = np.where(l_pred_test > 0.0, 1, -1)
error_test = 1 - np.mean(pred_test == y_test)
error_test_list.append(error_test)
results_list.append({"k": k, "Training Errors": error_train,
                    "TestErrors": error_test})

if k == 300:
    p_train_300 = l_pred_train
    pred_train_300 = pred_train

    p_test_300 = l_pred_test
    pred_train_300 = pred_test

    # Calculate ROC curve and ROC area for training
    train_fpr_300, train_tpr_300, _ = roc_curve(y_train, p_train_300)
    train_auc_300 = roc_auc_score(y_train, p_train_300)

    # Calculate ROC curve and ROC area for testing
    test_fpr_300, test_tpr_300, _ = roc_curve(y_test, p_test_300)
    test_auc_300 = roc_auc_score(y_test, p_test_300)

    # Plot ROC curves
    plt.figure(figsize=(10, 7))
    plt.plot(train_fpr_300, train_tpr_300,
             label=f'Training AUC (k=300): {train_auc_300:.2f}')
    plt.plot(test_fpr_300, test_tpr_300,
             label=f'Testing AUC (k=300): {test_auc_300:.2f}')
    plt.plot([0, 1], [0, 1], 'k--') # Dashed diagonal
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('ROC Curve for k=300')
    plt.legend(loc="lower right")
    plt.show()

results_df = pd.concat([pd.DataFrame(result, index=[0])
                        for result in results_list], ignore_index=True)

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(range(1, 601), loss_list)
plt.xlabel('Iteration')
plt.ylabel('Training Loss for K= 600')
plt.show()

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(k_values, error_train_list, label='Training')
plt.plot(k_values, error_test_list, label='Testing')
plt.xlabel('K Values')
plt.ylabel('Misclassification Error for Train and Test Data')
plt.legend()
plt.show()

return results_df

```

Load data

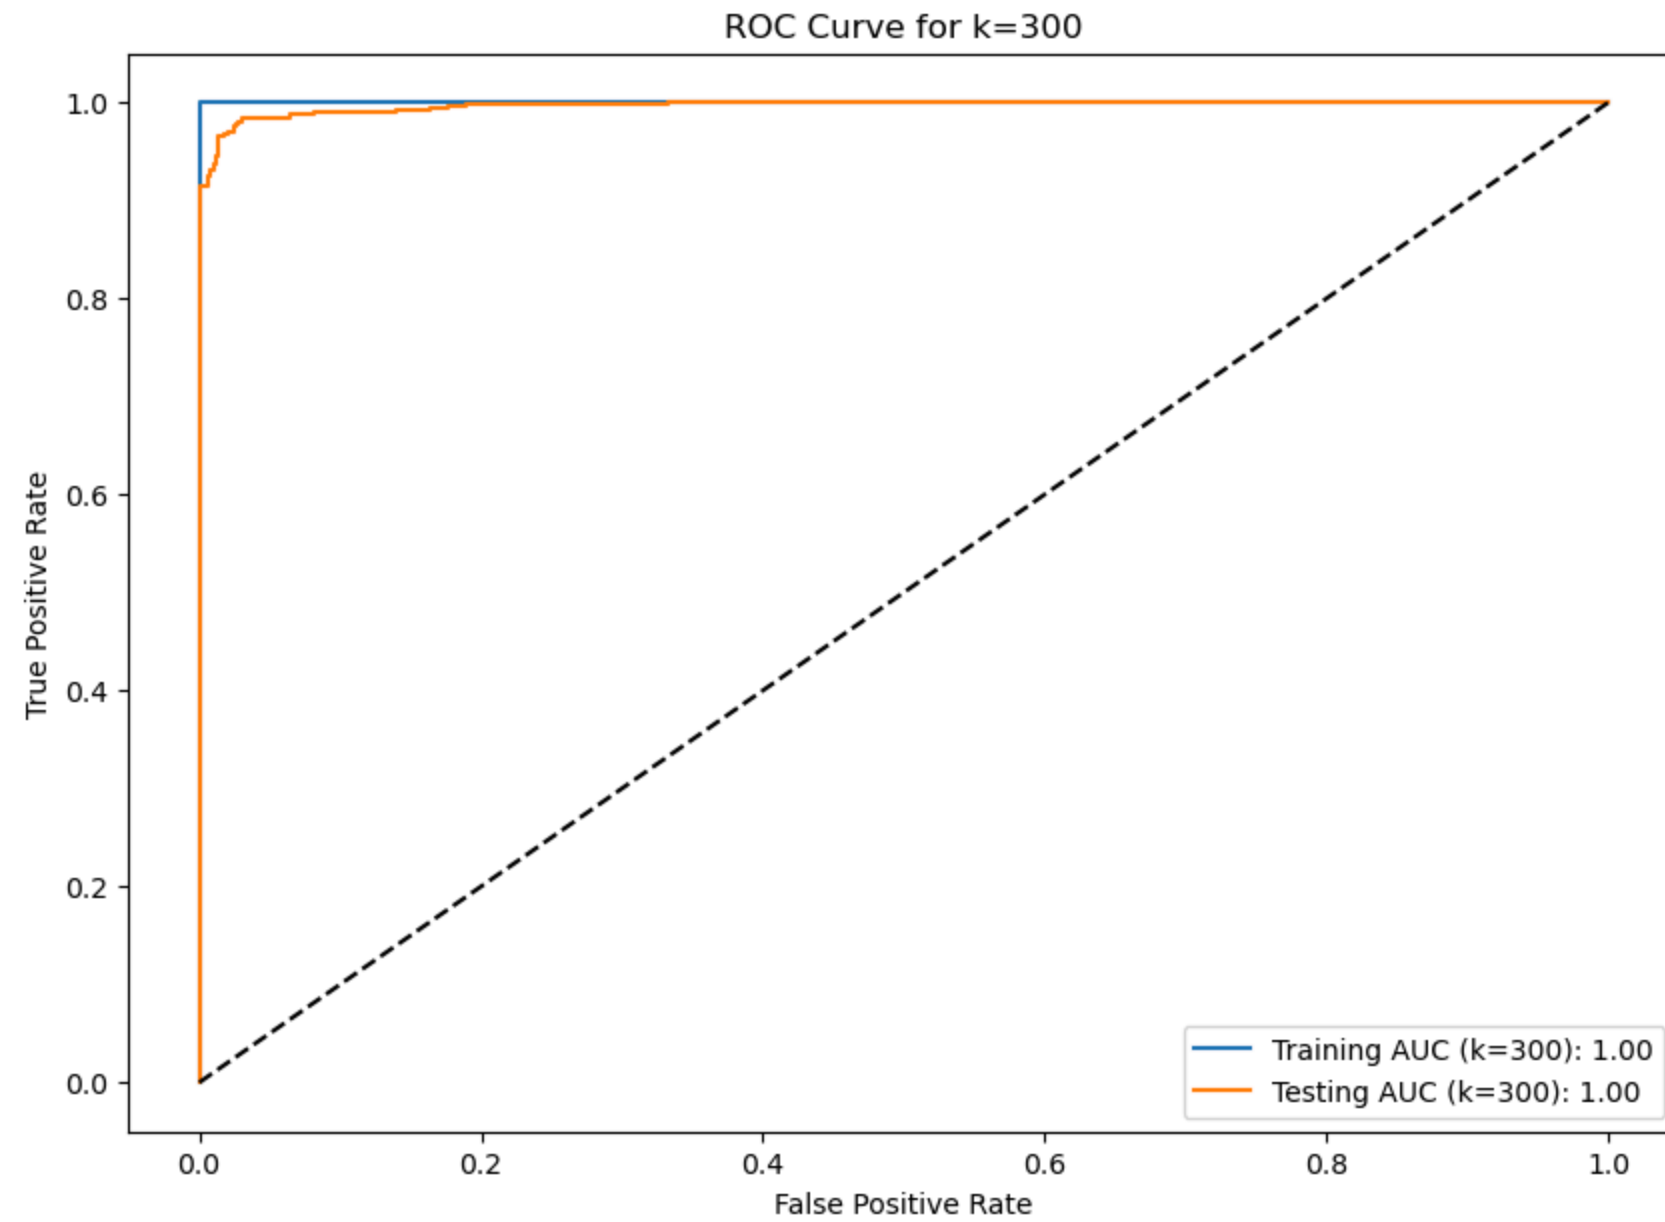
```
X_train_data = np.loadtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_4/Gisette/gisette_train.data")
y_train_labels = np.loadtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_4/Gisette/gisette_train.labels")



X_test_data = np.loadtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_4/Gisette/gisette_valid.data")
y_test_labels = np.loadtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_4/Gisette/gisette_valid.labels")

# Define values of k
k_values = [10, 30, 100, 300, 600]

# Call the logistic regression function
results_df = logistic_regression_logit(X_train_data, y_train_labels,
                                       X_test_data, y_test_labels, k_values)

print("Misclassification Error for Train and Test Data:")
print(results_df)
```





Misclassification Error for Train and Test Data:

	k	Training Errors	TestErrors
0	10	0.074667	0.087
1	30	0.033667	0.038
2	100	0.012500	0.030
3	300	0.000000	0.027
4	600	0.000000	0.020