Machine Learning Assignment

TISP variable selection

Question 1 (a)

In [1]:
```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import StandardScaler
import warnings
warnings.filterwarnings('ignore')
```

In [2]:
```
def TISP(X_Train, y_train, X_test, y_test, Lambda):
    X = X_Train
    Y = y_train
    X_test = X_test
    Y_test = y_test
    N, P = X.shape
    N_test = X_test.shape[0]
    W_old = np.zeros(P)
    train_errors_fr_iterations = []

    for i in range(100):
        gradient =  np.dot(X.T, Y - 1 / (1 + np.exp(-np.dot(X, W_old))))
        W_temp = W_old +  gradient
        W_new = W_temp * (np.abs(W_temp) > Lambda)
        W_old = W_new


        P_hat = 1 / (1 + np.exp(-np.dot(X, W_new)))
        Y_hat = 2 * (P_hat > 0.5) - 1
        Tr_table = pd.crosstab(Y, Y_hat)
        Tr_err = 1 - np.sum(np.diag(Tr_table)) / N
        train_errors_fr_iterations.append(Tr_err)

    var = np.sum(W_new != 0)
    P_hat = 1 / (1 + np.exp(-np.dot(X, W_new)))
    Y_hat = 2 * (P_hat > 0.5) - 1
    Tr_table = pd.crosstab(Y, Y_hat)
    Tr_err = 1 - np.sum(np.diag(Tr_table)) / N


    P_hat_test = 1 / (1 + np.exp(-np.dot(X_test, W_new)))
    Y_hat_test = 2 * (P_hat_test > 0.5) - 1
    Ts_table = pd.crosstab(Y_test, Y_hat_test)
    Ts_err = 1 - np.sum(np.diag(Ts_table)) / N_test

    return {'W_hat': W_new,
            'Train_error': Tr_err,
            'Test_error': Ts_err,
            'Selected_Features': var,
            'Given_Lambda': Lambda,
            'Train_errors': train_errors_fr_iterations}
```

In [3]:
```
X_train = np.loadtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_4/Gisette/gisette_train.data")
y_train = np.loadtxt("//Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_4/Gisette/gisette_train.labels")

X_test = np.loadtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_4/Gisette/gisette_valid.data")
y_test = np.loadtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_4/Gisette/gisette_valid.labels")

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

y_train = (y_train + 1) // 2
y_test = (y_test + 1) // 2
```

In [4]:
```
lambda_values =[1007.900,880.099,985,262.0, 201.500]
M = len(lambda_values)
TISP_var = np.zeros(M)
TISP_Tr_err = np.zeros(M)
TISP_Ts_err = np.zeros(M)
TISP_weights = []
train_errors_vs_iteration = []

for i in range(M):
    TISP_result = TISP(X_train, y_train, X_test, y_test, Lambda=lambda_values[i])
    TISP_var[i] = TISP_result['Selected_Features']
    TISP_Tr_err[i] = TISP_result['Train_error']
    TISP_Ts_err[i] = TISP_result['Test_error']
    TISP_weights.append(TISP_result['W_hat'])
    train_errors_vs_iteration.append(TISP_result['Train_errors'])

plt.plot(range(1, 101), train_errors_vs_iteration[2],color='darkorange')
plt.xlabel("No of Iterations")
plt.ylabel("Misclassification Error")
plt.title("Train Misclassification Error v/s Iteration (100 Features)")
plt.show()
plt.plot(TISP_var, TISP_Tr_err, label="Train Error")
plt.plot(TISP_var, TISP_Ts_err, label="Test Error")
plt.xlabel("Number of Selected Features")
plt.ylabel("Misclassification Error")
plt.title("Train and Test Misclassification Error v/s Number of Selected Features")
plt.legend()
plt.show()

results = pd.DataFrame({ "Given Lambda": lambda_values,
                         "Selected Features": TISP_var,
                         "Train Error": TISP_Tr_err,
                         "Test Error": TISP_Ts_err
})

print(results)

fpr_test, tpr_test, _ = roc_curve(y_test, 1
                                  / (1 + np.exp(-np.dot(X_test, TISP_weights[2]))))
roc_auc_test = auc(fpr_test, tpr_test)

fpr_train, tpr_train, _ = roc_curve(y_train, 1
                                    / (1 + np.exp(-np.dot(X_train, TISP_weights[2]))))
roc_auc_train = auc(fpr_train, tpr_train)

plt.figure()
plt.plot(fpr_train, tpr_train, color='darkorange', lw=2,
         label=f'Train ROC curve (area = {roc_auc_train:.2f})')
plt.plot(fpr_test, tpr_test, color='blue', lw=2,
         label=f'Test ROC curve (area = {roc_auc_test:.2f})')
```

```
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()
```

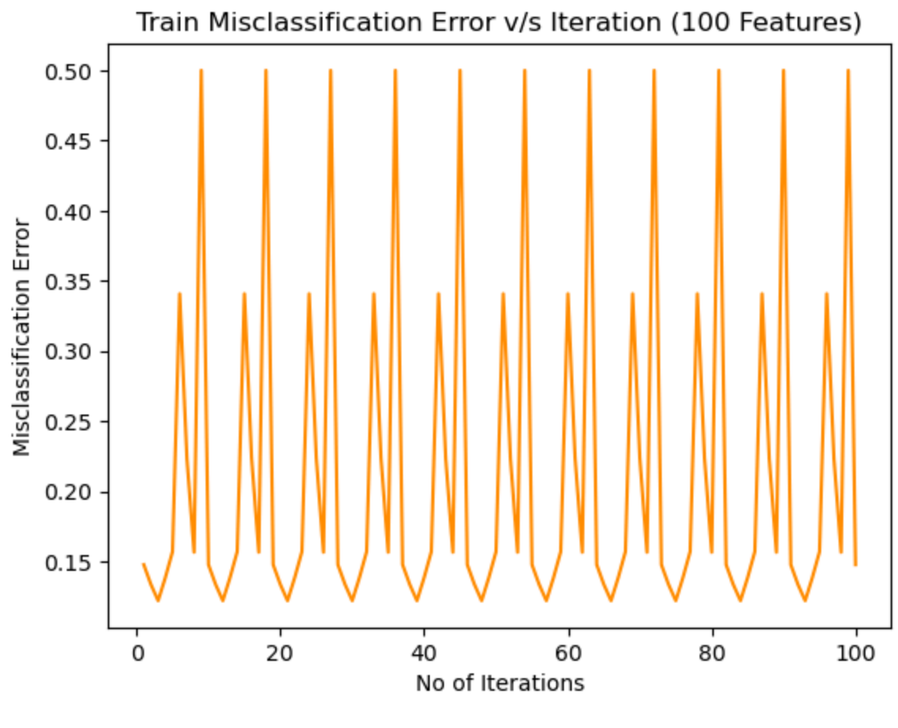

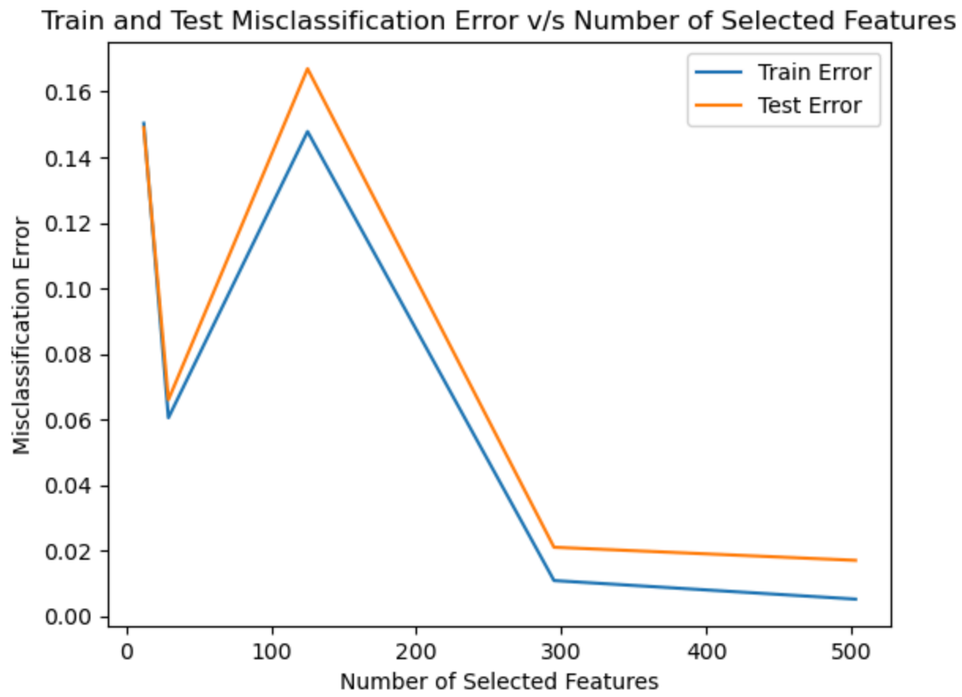

	Given Lambda	Selected Features	Train Error	Test Error
0	1007.900	12.0	0.150333	0.149
1	880.099	29.0	0.060500	0.066
2	985.000	125.0	0.147833	0.167
3	262.000	295.0	0.010833	0.021
4	201.500	503.0	0.005167	0.017

Question 1 (b)

In [5]:
```
from sklearn.preprocessing import StandardScaler, MinMaxScaler
```

In [6]:
```
X_train = np.loadtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_4/MADELON/madelon_train.data")
y_train = np.loadtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_4/MADELON/madelon_train.labels")

X_test = np.loadtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_4/MADELON/madelon_valid.data")
y_test = np.loadtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_4/MADELON/madelon_valid.labels")
```

In [7]:
```
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)


y_train = (y_train + 1) // 2
y_test = (y_test + 1) // 2
```

In [8]:
```
lambda_values =[160.50,119.0,83.19,42.30,1]
M = len(lambda_values)
TISP_var = np.zeros(M)
TISP_Tr_err = np.zeros(M)
TISP_Ts_err = np.zeros(M)
TISP_weights = []
train_errors_vs_iteration = []

for i in range(M):
    TISP_result = TISP(X_train, y_train, X_test, y_test, Lambda=lambda_values[i])
```

```
    TISP_var[i] = TISP_result['Selected_Features']
    TISP_Tr_err[i] = TISP_result['Train_error']
    TISP_Ts_err[i] = TISP_result['Test_error']
    TISP_weights.append(TISP_result['W_hat'])
    train_errors_vs_iteration.append(TISP_result['Train_errors'])

plt.plot(range(1, 101), train_errors_vs_iteration[2],color='darkorange')
plt.xlabel("No of Iterations")
plt.ylabel("Misclassification Error")
plt.title("Train Misclassification Error v/s Iteration (100 Features)")
plt.show()
plt.plot(TISP_var, TISP_Tr_err, label="Train Error")
plt.plot(TISP_var, TISP_Ts_err, label="Test Error")
plt.xlabel("Number of Selected Features")
plt.ylabel("Misclassification Error")
plt.title("Train and Test Misclassification Error v/s Number of Selected Features")
plt.legend()
plt.show()

results = pd.DataFrame({ "Given Lambda": lambda_values,
                         "Selected Features": TISP_var,
                         "Train Error": TISP_Tr_err,
                         "Test Error": TISP_Ts_err
})

print(results)

fpr_test, tpr_test, _ = roc_curve(y_test, 1
                                  / (1 + np.exp(-np.dot(X_test, TISP_weights[2]))))
roc_auc_test = auc(fpr_test, tpr_test)

fpr_train, tpr_train, _ = roc_curve(y_train, 1
                                    / (1 + np.exp(-np.dot(X_train, TISP_weights[2]))))
roc_auc_train = auc(fpr_train, tpr_train)

plt.figure()
plt.plot(fpr_train, tpr_train, color='darkorange', lw=2,
         label=f'Train ROC curve (area = {roc_auc_train:.2f})')
plt.plot(fpr_test, tpr_test, color='blue', lw=2,
         label=f'Test ROC curve (area = {roc_auc_test:.2f})')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()
```

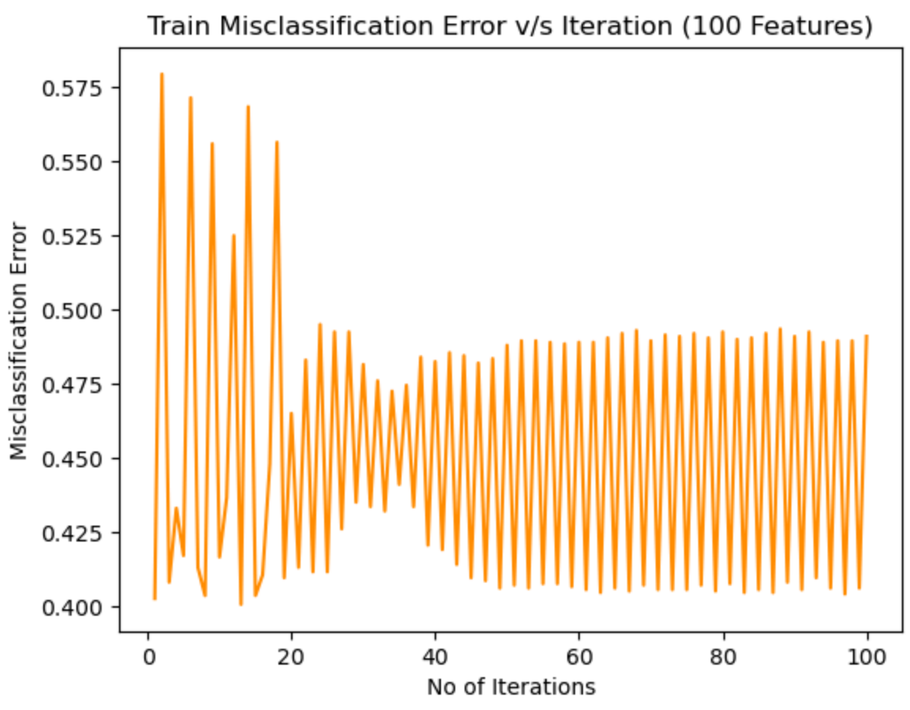

Train and Test Misclassification Error v/s Number of Selected Features

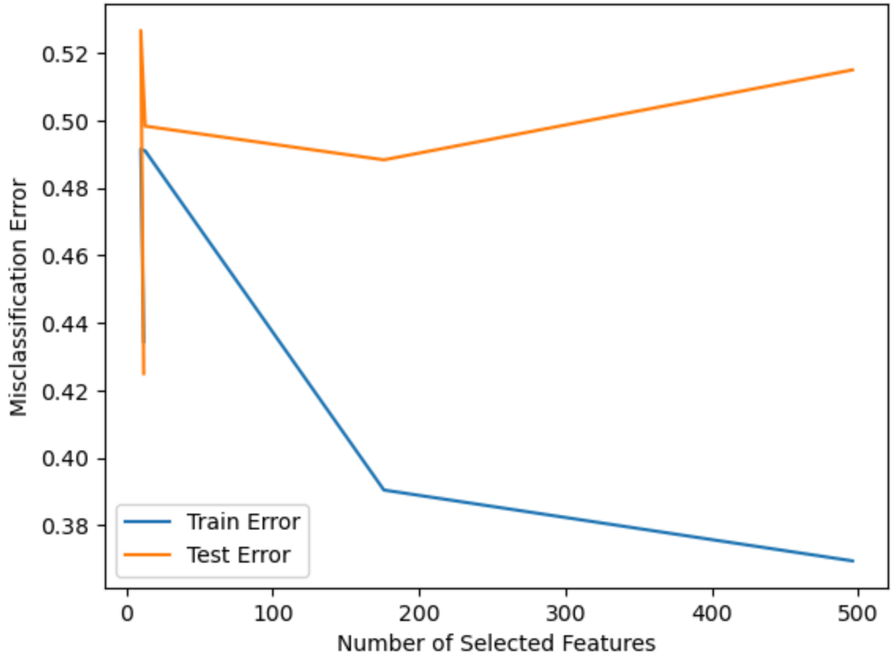

	Given Lambda	Selected Features	Train Error	Test Error
0	160.50	12.0	0.4345	0.425000
1	119.00	10.0	0.4915	0.526667
2	83.19	13.0	0.4910	0.498333
3	42.30	176.0	0.3905	0.488333
4	1.00	496.0	0.3695	0.515000

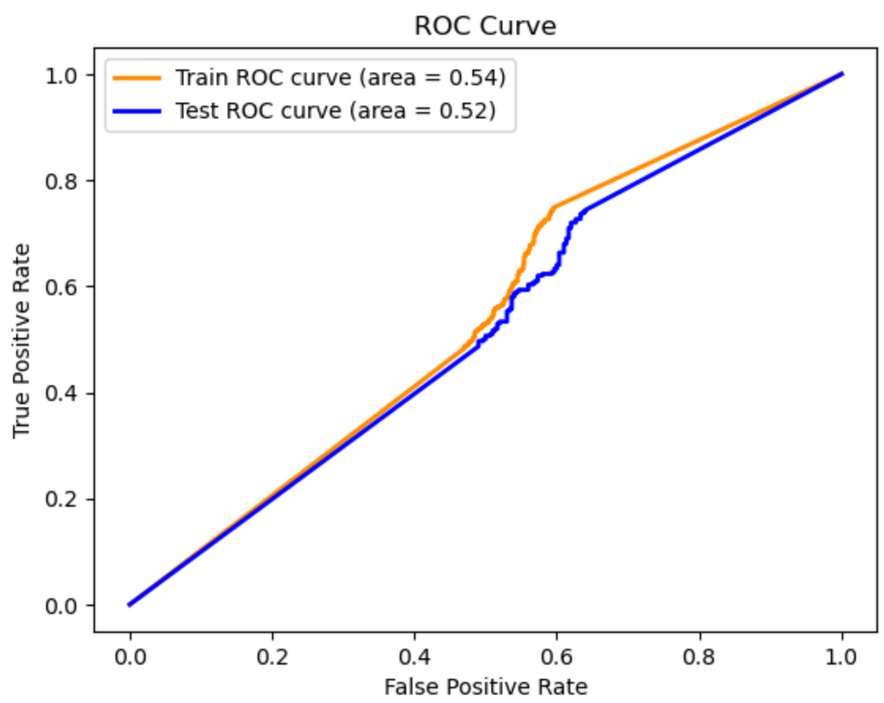

Question 1 (c)

In [9]:
```
X_train = np.genfromtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_3/dexter/dexter_train.csv", delimiter=',')
y_train = np.genfromtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_3/dexter/dexter_train.labels")

X_test = np.genfromtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_3/dexter/dexter_valid.csv", delimiter=',')
y_test = np.genfromtxt("/Users/gaganullas19/Documents/Spring2024/AppliedMachineLearning/Homework_3/dexter/dexter_valid.labels")

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Converting -1 to 0 and keeping 1 as 1
y_train = (y_train + 1) // 2
y_test = (y_test + 1) // 2
```

In [10]:
```
lambda_values =[32.001, 25,20.6,16.999,15.135]
M = len(lambda_values)
TISP_var = np.zeros(M)
TISP_Tr_err = np.zeros(M)
TISP_Ts_err = np.zeros(M)
TISP_weights = []
train_errors_vs_iteration = []

for i in range(M):
    TISP_result = TISP(X_train, y_train, X_test, y_test, Lambda=lambda_values[i])
    TISP_var[i] = TISP_result['Selected_Features']
    TISP_Tr_err[i] = TISP_result['Train_error']
    TISP_Ts_err[i] = TISP_result['Test_error']
    TISP_weights.append(TISP_result['W_hat'])
    train_errors_vs_iteration.append(TISP_result['Train_errors'])

plt.plot(range(1, 101), train_errors_vs_iteration[2],color='darkorange')
plt.xlabel("No of Iterations")
plt.ylabel("Misclassification Error")
```

```
plt.title("Train Misclassification Error v/s Iteration (100 Features)")
plt.show()
plt.plot(TISP_var, TISP_Tr_err, label="Train Error")
plt.plot(TISP_var, TISP_Ts_err, label="Test Error")
plt.xlabel("Number of Selected Features")
plt.ylabel("Misclassification Error")
plt.title("Train and Test Misclassification Error v/s Number of Selected Features")
plt.legend()
plt.show()

results = pd.DataFrame({ "Given Lambda": lambda_values,
                         "Selected Features": TISP_var,
                         "Train Error": TISP_Tr_err,
                         "Test Error": TISP_Ts_err
})

print(results)

fpr_test, tpr_test, _ = roc_curve(y_test, 1 /
                                  (1 + np.exp(-np.dot(X_test, TISP_weights[2]))))
roc_auc_test = auc(fpr_test, tpr_test)

fpr_train, tpr_train, _ = roc_curve(y_train, 1 /
                                    (1 + np.exp(-np.dot(X_train, TISP_weights[2]))))
roc_auc_train = auc(fpr_train, tpr_train)

plt.figure()
plt.plot(fpr_train, tpr_train, color='darkorange', lw=2,
         label=f'Train ROC curve (area = {roc_auc_train:.2f})')
plt.plot(fpr_test, tpr_test, color='blue', lw=2,
         label=f'Test ROC curve (area = {roc_auc_test:.2f})')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.show()
```

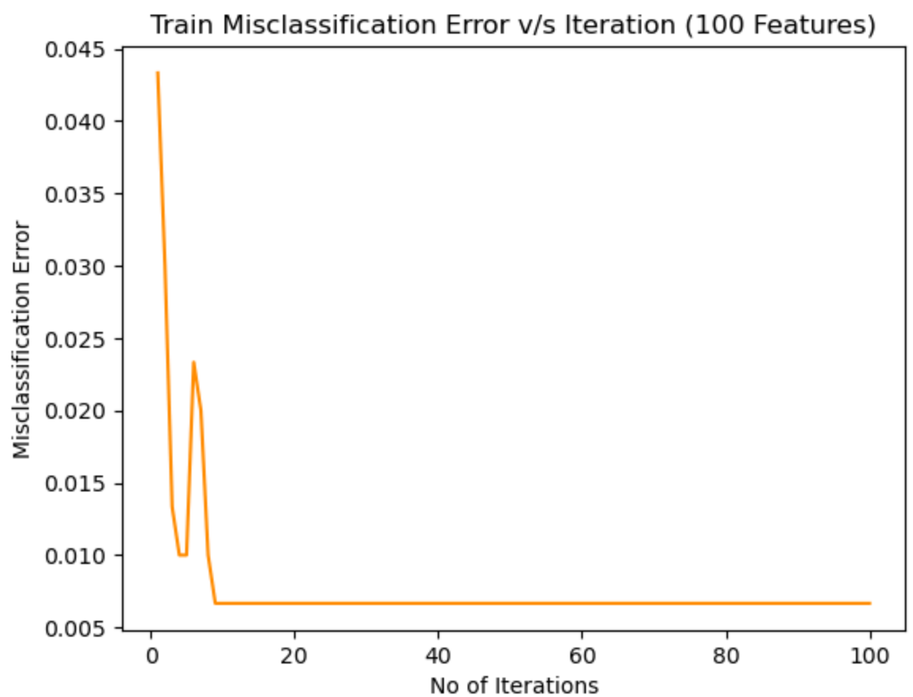

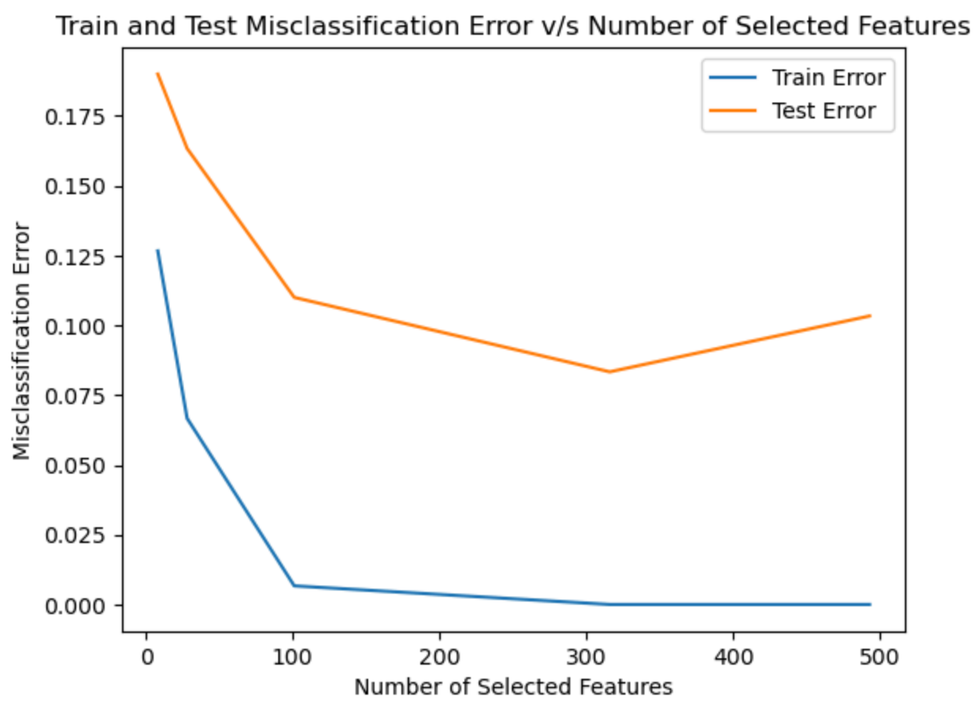

	Given Lambda	Selected Features	Train Error	Test Error
0	32.001	8.0	0.126667	0.190000
1	25.000	28.0	0.066667	0.163333
2	20.600	101.0	0.006667	0.110000
3	16.999	316.0	0.000000	0.083333
4	15.135	493.0	0.000000	0.103333